SaliencySlider: Enhancing AI Transparency through Interactive Visual Interpretation of CNNs

Matheus Kunzler Maldaner[1], Kian Ambrose[1],
Kristian O'Connor[1], Alexa Certo[1]

[1]University of Florida, United States

Abstract

The increasing integration of Artificial Intelligence (AI) systems in critical sectors necessitates not only high accuracy but also a clear understanding of their decision-making processes. To address the need for transparency in AI, we introduce the SaliencySlider, an interactive web tool that utilizes Gradient-weighted Class Activation Mapping (Grad-CAM) to visually explain the decisions of convolutional neural networks. Users can dynamically adjust a slider to visualize how different regions of an input image influence the AI's classification decisions, thus making the model's internal reasoning accessible and understandable. The tool supports the auditing of AI decisions by highlighting saliency in real-time, fostering trust and facilitating deeper insights into model behavior. We discuss the development and deployment of SaliencySlider, its underlying technology, and its potential impact on promoting explainable AI in applications where understanding model reasoning is crucial. Preliminary user feedback underscores its effectiveness and potential as an educational and diagnostic tool in various domains.

Keywords: Explainable AI, Grad-CAM, Visual Interpretability, Saliency Mapping, CNNs.

1. Introduction

In the era of Artificial Intelligence (AI), the adoption of complex machine learning models, particularly deep learning networks, has led to significant advancements in various sectors including healthcare, finance, and autonomous systems. However, the "black-box" nature of these models poses challenges in understanding their decision-making processes, which can hinder trust and accountability, especially in critical applications. This concern has catalyzed the development of explainable AI (XAI), which aims to make AI decisions transparent, understandable, and trustworthy.

Among various XAI techniques, saliency mapping stands out as an intuitive method to visualize the influence of different parts of input data on the output of neural networks. Our project, SaliencySlider, leverages Gradient-weighted Class Activation Mapping (Grad-CAM), a sophisticated saliency mapping technique, to develop an interactive tool that enhances the transparency of convolutional neural networks (CNNs). By enabling users to dynamically adjust the visualization of saliency regions through a slider interface, the tool allows for an insightful exploration into how specific image features affect model predictions.

The motivation behind SaliencySlider was to bridge the gap between AI functionalities and user understanding, thereby enhancing the interpretability of complex models. It was conceived with the objective of not only providing a means to visualize model decisions but also to educate and empower users, allowing them to critically evaluate AI outputs. The final implementation of the tool as a web application ensures accessibility and ease of use, making it a practical resource for both novices and experts interested in the workings of AI.

Our tool does not stop at providing visual explanations; it also serves as a platform for research and development, offering insights that could lead to more robust and fair AI systems. This paper describes the development and deployment of SaliencySlider, discusses its underlying methods, and evaluates its impact in advancing the transparency of AI models.

2. Methods

System Architecture and Deployment

SaliencySlider integrates TensorFlow and Keras for leveraging the VGG19 model, a robust convolutional neural network known for its effective image classification. The application is hosted on pythonanywhere to ensure global accessibility and is constructed using a frontend of HTML, CSS, and JavaScript for user interaction, with Django as the backend framework to handle server-side logic.

Image Processing and Model Prediction

Images uploaded by users are resized to 224x224 pixels and preprocessed to align with ImageNet standards, involving normalization and color adjustment. The prepared images are fed into VGG19, which outputs predictions and confidence scores. The predictions are utilized both for direct user feedback and as inputs into the Grad-CAM process for saliency mapping.

Implementation of Grad-CAM

Grad-CAM is applied to visually interpret the CNN's decision-making. The process includes:

1. Forward Propagation:
The image passes through the CNN, yielding class scores at the output layer. This step is mathematically represented by the function $f(x)$ over the input image x, culminating in class scores.

2. Target Class Identification:
The class for which the saliency map is generated is identified, typically the top prediction.

3. Gradient Calculation:
Gradients of the target class score, y^c, with respect to the feature maps, A^k, of a chosen convolutional layer are calculated:

$$\frac{\partial y^c}{\partial A^k}$$

4. Pooling of Gradients:
These gradients are pooled to derive the weights α_k^c for the feature maps:

$$\alpha_k^c = \frac{1}{Z} \sum_i \sum_j \frac{\partial y^c}{\partial A_{ij}^k}$$

where Z represents the total number of elements in feature map A^k, indexing the spatial dimensions i, j.

5. Heatmap Generation:

A weighted combination of feature maps is computed and passed through a ReLU function to ensure non-negative emphasis, forming the final visual saliency map:

$$L^c_{\text{Grad-CAM}} = \text{ReLU}\left(\sum_k \alpha^c_k A^k\right)$$

Interactive Saliency Visualization

The resulting heatmap is overlaid on the original image with adjustable opacity controlled by a slider, allowing users to visualize the impact of different image regions on the model's prediction dynamically. This interaction not only clarifies which features are most influential but also provides an educational insight into neural network processing.

User Interface and Experience

The design of the user interface prioritizes simplicity and educational utility, facilitating ease of use and accessibility for a broad audience. Detailed tooltips and guides enhance the user's ability to interact effectively with the SaliencySlider, promoting a deeper understanding of AI transparency.

3. Discussion

The development of SaliencySlider addresses critical needs in the field of explainable AI by providing a practical tool for visualizing how neural networks process visual information. The ability to dynamically explore the influence of different image regions on model predictions not only enhances user understanding but also aids in the diagnostic process, identifying potential biases or errors in the AI's decision-making. However, initial feedback has indicated some areas for improvement, particularly in terms of interface responsiveness and feature robustness. Addressing these concerns will be crucial as we continue to refine the tool.

4. Future Work

Future developments for SaliencySlider will focus on expanding its capabilities to include additional models and interpretability techniques, enhancing the user interface, and optimizing performance for a broader range of devices. We also plan to conduct a structured user study to collect detailed feedback and usage data, which will inform further improvements. Additionally, incorporating user-generated content and community features may enhance the tool's educational and collaborative potentials.

5. Conclusion

SaliencySlider represents a significant step forward in making AI systems more transparent and understandable. By providing a user-friendly platform for exploring AI decision-making visually, it bridges the gap between complex machine learning technology and practical, everyday applications. As the tool evolves, it holds the promise of significantly impacting how people interact with and trust AI systems, promoting greater adoption and more informed usage of AI across various sectors.